

Assignment 01 - Intensity Transformations and Neighborhood Filtering

EN3160: Image Processing And Machine Vision

Index: 210349N - Madhushan I.D.



GitHub Link for Code: [Link to my GitHub repository](#)

Question 1: Intensity Transformation

Resulting Image:



Figure 1: Result of the intensity transformation applied on the image.

Mid-range intensities(50-150) are enhanced, while the low (0-50) and high (150-255) intensities remain unchanged. Results in brighter mid-tone areas such as some parts of the right cheek and forehead, while darker and lighter regions such as shadows and highlights stay the same.

Question 2: Enhancing White and Gray Matter

Results:

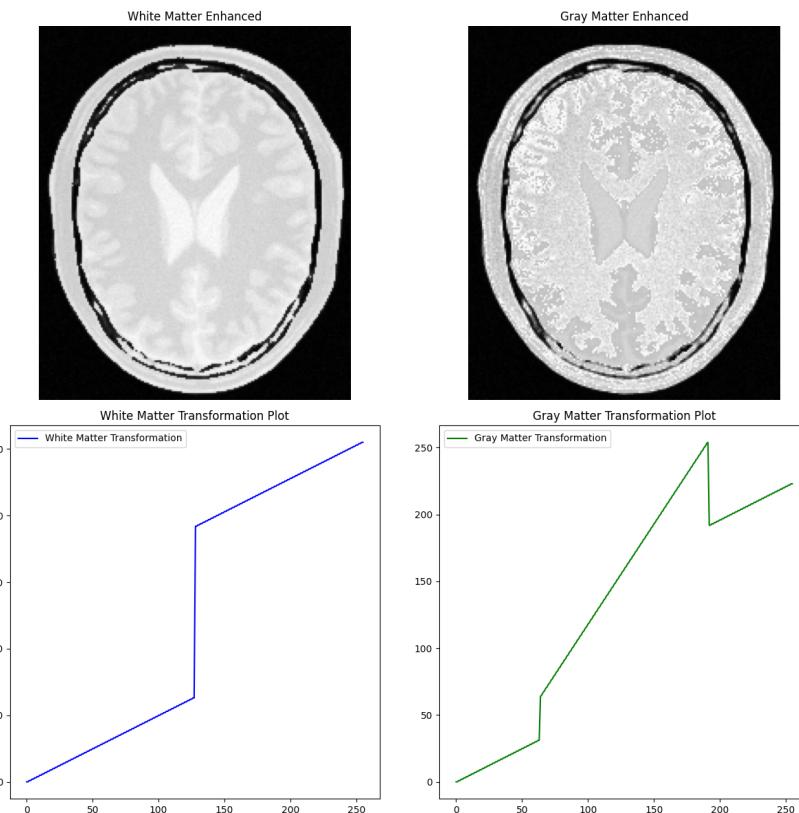


Figure 2: Enhanced White matter (left) and Gray matter (right).

White matter transformation enhances high intensities, while gray matter transformation enhances mid intensities.

Question 3: Gamma Correction

(a) $\gamma = 0.5$

Code:

```

1 L, a, b = cv2.split(image_lab) # Separate the L, a, b channels
2 def adjust_gamma(image, gamma=1.0): # Function to apply gamma correction
3     table = np.array([(i / 255.0) ** gamma) * 255 for i in np.arange(0, 256)]).astype("uint8")
4     return cv2.LUT(image, table)
5
6 gamma_corrected_L = adjust_gamma(L, gamma=0.5) # Apply gamma correction with gamma = 0.5

```

(b) Results:



Figure 3: Original image(left) and gamma-corrected image(right).

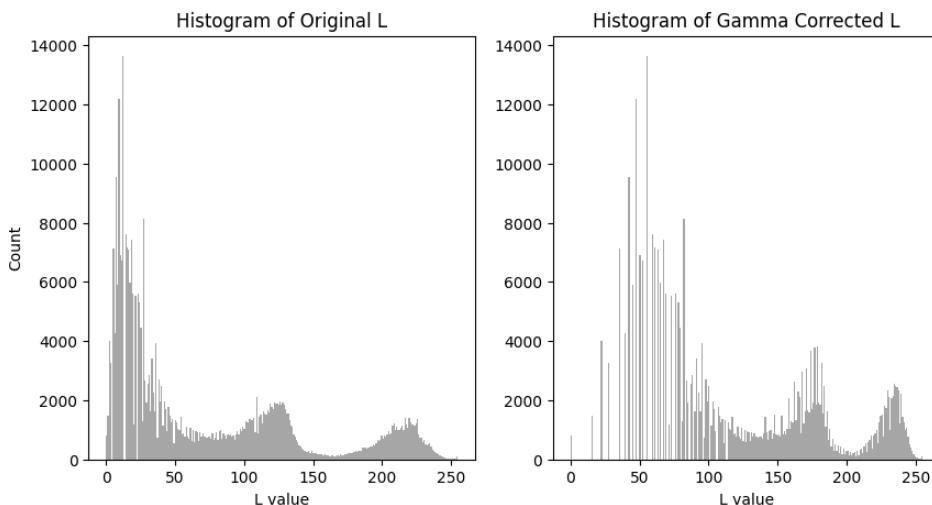


Figure 4: Histograms of the original(left) and gamma-corrected(right) images.

Question 4: Enhancing the Vibrance of a Photograph

Part (a): Splitting the Image

```

1 hue, saturation, value = cv2.split(image_hsv) # Split into hue, saturation, and value components

```

Part (b): Applying the Intensity Transformation

```

1 def intensity_transform(x, alpha=0.5, sigma=70): # Define the transformation function
2     return np.minimum(x + alpha * (128 * np.exp(-((x-128)**2) / (2*sigma**2))), 255)
3 saturation_transformed = intensity_transform(saturation, alpha=0.6).astype(np.uint8)

```

Part (c): Adjusting the Parameter α

The optimal value found was $\alpha = 0.5$.

```

1 from ipywidgets import interact, FloatSlider
2 def adjust_saturation_plane(alpha):
3     saturation_transformed = intensity_transform(saturation, alpha=alpha).astype(np.uint8)
4     plt.figure(figsize=(6, 4))
5     plt.imshow(saturation_transformed, cmap='gray')
6     plt.colorbar()
7     plt.title(f'Transformed Saturation Plane (alpha={alpha:.2f})')
8     plt.axis('off')
9     plt.show()
10 # Create an interactive slider
11 interact(adjust_saturation_plane, alpha=FloatSlider(value=0.5, min=0, max=1, step=0.05, description='Alpha'))

```

Part (d): Recombining the Planes

```

1 image_hsv_transformed = cv2.merge([hue, saturation_transformed, value]) # Reconstruct the image
2 image_transformed = cv2.cvtColor(image_hsv_transformed, cv2.COLOR_HSV2BGR)

```

Part (e): Displaying results



Figure 5: Original



Figure 6: Vibrance-enhanced

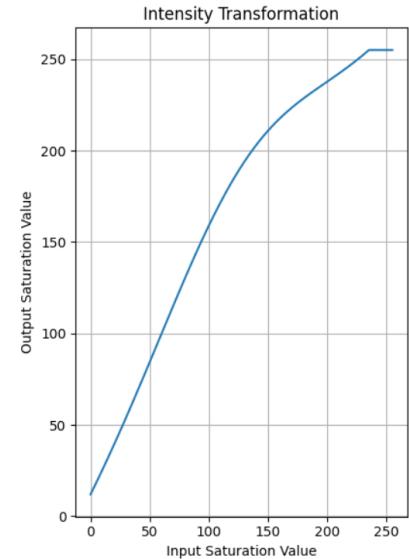


Figure 7: Intensity transformation

Question 5: Histogram Equalization

Function to carry out histogram equalization:

```

1 def calculate_histogram(image):
2     histogram = np.zeros(256, dtype=int)
3     for pixel in image.flatten():
4         histogram[pixel] += 1
5     return histogram
6 def equalize_histogram(image):
7     hist = calculate_histogram(image) # Calculate the histogram
8     cdf = np.cumsum(hist) # Calculate the cumulative distribution function (CDF)
9     cdf_normalized = (cdf) * 255 / cdf.max() # Normalize the CDF
10    cdf_normalized = cdf_normalized.astype(np.uint8)
11    equalized_image = cdf_normalized[image] # Use the normalized CDF to map the original image pixels to
12    return equalized_image, hist, calculate_histogram(equalized_image)

```

Result:

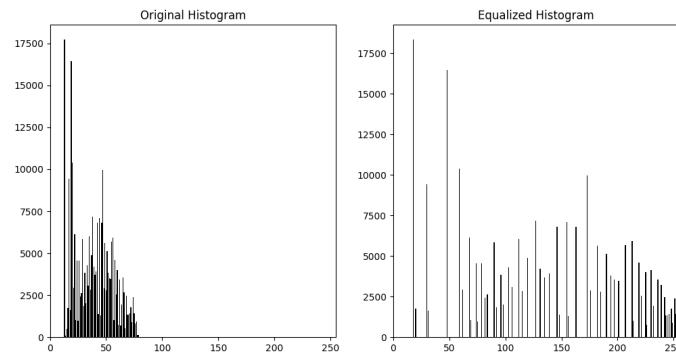


Figure 8: Histograms before and after equalization.

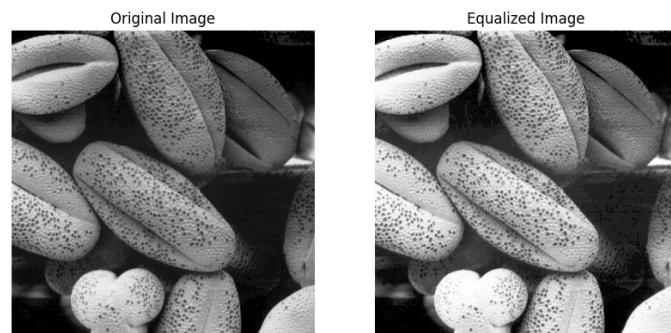


Figure 9: Comparison between the original image (left) and the histogram-equalized image (right).

Question 6: Foreground Histogram Equalization

Part (a): Splitting the Image into HSV Planes

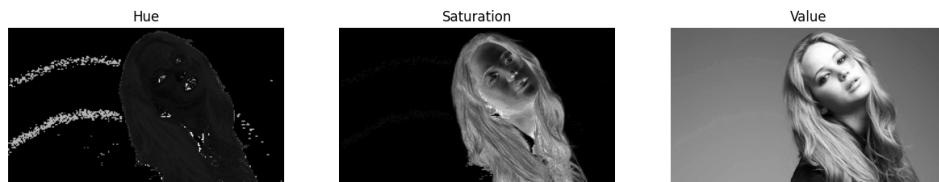


Figure 10: Hue, Saturation, and Value components of the image.

Part (b): Thresholding to Extract the Foreground

The saturation component was used to threshold and extract the foreground mask.

```
1 _, mask = cv2.threshold(saturation, 11, 255, cv2.THRESH_BINARY) #saturation plane for extracting the foreg
```



Figure 11: Binary mask of the foreground and foreground.

Part (c): Obtain the foreground using cv.bitwise_and and compute the histogram

```
1 foreground = cv2.bitwise_and(image_rgb, image_rgb, mask=mask) # Extract the foreground using the mask
2 histogram = calculate_histogram(foreground) # Calculate histogram of the foreground
```

Part (d): Cumulative Sum

```
1 cdf = np.cumsum(histogram) # Calculate the cumulative sum of the histogram
```

Part (e): histogram-equalize the foreground

```
1 cdf_normalized = (cdf) * histogram.max() / (cdf.max()) # Normalize the CDF
2 cdf_normalized = cdf_normalized.astype(np.uint8)
3 equalized_value = cdf_normalized[foreground] # Map the value plane
```

Part (f): Extract the background and add with the histogram equalized foreground

```
1 background_image = cv2.bitwise_and(image_rgb, image_rgb, mask=~mask) # extract background from original image
2 final_image = cv2.add(background_image, foreground) # Combine background and equalized foreground
```

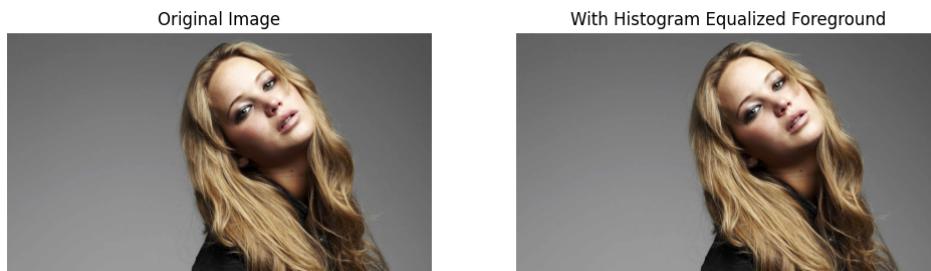


Figure 12: original image, and the result with the histogram equalized foreground.

Question 7: Sobel Filtering**Part (a): Using the Existing Filter**

```
1 sobel_x = np.array([[1, 0, -1], [2, 0, -2], [1, 0, -1]], dtype="int")
2 sobel_y = np.array([[1, 2, 1], [0, 0, 0], [-1, -2, -1]], dtype="int")
3 sobel_filtered_x = cv2.filter2D(image, -1, sobel_x)
4 sobel_filtered_y = cv2.filter2D(image, -1, sobel_y)
```



Figure 13: Sobel filtering in the X direction (left) and Y direction (right).

Part (b): Writing Custom Code for Sobel Filtering

```
1 def apply_sobel_filter(image, kernel):
2     rows, cols = image.shape
3     padded_image = np.pad(image, 1, mode='constant') # Pad the image
4     output = np.zeros_like(image) # Initialize the output image
5     for i in range(1, rows-1): # Perform convolution
6         for j in range(1, cols-1):
7             roi = padded_image[i-1:i+2, j-1:j+2] # Extract the region of interest
```

```

8     output[i, j] = np.sum(roi * kernel) # Perform element-wise multiplication
9
10    return output
11
12 sobel_x_output = apply_sobel_filter(image, sobel_x) # Apply Sobel filter
13 sobel_y_output = apply_sobel_filter(image, sobel_y)

```

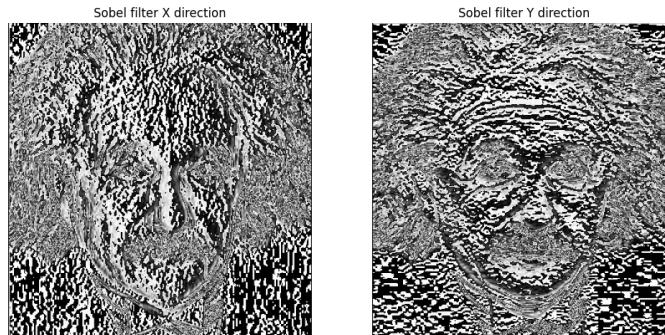


Figure 14: Custom Sobel filtering in the X direction (left) and Y direction (right).

Part (c): Using the Property of Separable Filters

```

1 # Define the 1D Sobel kernels
2 sobel_vertical = np.array([[1], [2], [1]]) # Vertical kernel
3 sobel_horizontal = np.array([[1, 0, -1]]) # Horizontal kernel
4 # Apply the vertical Sobel filter
5 gradient_x = cv2.filter2D(image_gray, cv2.CV_64F, sobel_horizontal) # Using horizontal kernel
6 # Now apply the vertical kernel to the result
7 sobel_filtered_image = cv2.filter2D(gradient_x, cv2.CV_64F, sobel_vertical)
8 # Normalize the output image to the range [0, 255]
9 sobel_filtered_image=cv2.normalize(sobel_filtered_image, None, 0, 255, cv2.NORM_MINMAX).astype(np.uint8)
10

```

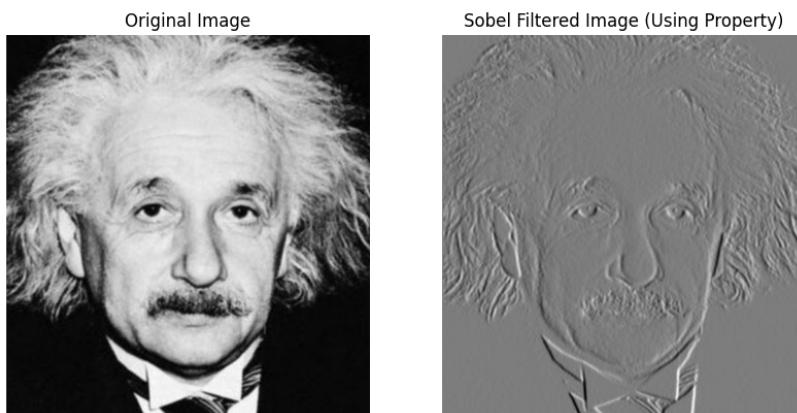


Figure 15: Result of applying separable filters for Sobel operation.

Question 8: Image Zooming

Part (a): Nearest-Neighbor Interpolation

```

1 def nearest_neighbor_zoom(image, scale_factor):
2     src_height, src_width = image.shape[:2]
3     dst_height, dst_width = int(src_height * scale_factor), int(src_width * scale_factor)
4     zoomed_image = np.zeros((dst_height, dst_width, image.shape[2]), dtype=image.dtype)
5     for i in range(dst_height):
6         for j in range(dst_width):
7             src_x = int(i / scale_factor)

```

```

8         src_y = int(j / scale_factor)
9         zoomed_image[i, j] = image[src_x, src_y]
10    return zoomed_image

```

Part (b): Bilinear Interpolation

```

1 def bilinear_interpolation_zoom(image, scale_factor):
2     src_height, src_width = image.shape[:2]
3     dst_height, dst_width = int(src_height * scale_factor), int(src_width * scale_factor)
4     zoomed_image = np.zeros((dst_height, dst_width, image.shape[2]), dtype=image.dtype)
5     for i in range(dst_height):
6         for j in range(dst_width):
7             x = i / scale_factor
8             y = j / scale_factor
9             x0 = int(x)
10            y0 = int(y)
11            x1 = min(x0 + 1, src_height - 1)
12            y1 = min(y0 + 1, src_width - 1)
13            a = x - x0
14            b = y - y0
15            top = (1 - a) * image[x0, y0] + a * image[x0, y1]
16            bottom = (1 - a) * image[x1, y0] + a * image[x1, y1]
17            zoomed_image[i, j] = (1 - b) * top + b * bottom
18    return zoomed_image.astype(image.dtype)

```

Testing and Comparison

```

1 def compute_ssd(original, scaled):
2     return np.mean((original - scaled) ** 2)
3 original_paths = ["im01.png", "im02.png"]
4 small_paths = ["im01small.png", "im02small.png"]
5 scale_factor = 4
6 for i in range(len(original_paths)):
7     original_image = cv2.imread(original_paths[i], cv2.IMREAD_COLOR)
8     small_image = cv2.imread(small_paths[i], cv2.IMREAD_COLOR)
9     nn_zoomed = nearest_neighbor_zoom(small_image, scale_factor)
10    bi_zoomed = bilinear_interpolation_zoom(small_image, scale_factor)
11    ssd_nn = compute_ssd(original_image, nn_zoomed)
12    ssd_bi = compute_ssd(original_image, bi_zoomed)
13    print(f"SSD with Nearest Neighbor for Image {i+1}: {ssd_nn}")
14    print(f"SSD with Bilinear Interpolation for Image {i+1}: {ssd_bi}")

```

Results

Image & Method	SSD with Nearest Neighbor	SSD with Bilinear Interpolation
Image 1	31.284316486625514	43.3944849537037
Image 2	11.902013310185184	18.214121238425925

For both images, the nearest neighbor interpolation method resulted in lower SSD values compared to bilinear interpolation. This suggests that, nearest neighbor interpolation more closely approximates the original images.

Question 9: Image Enhancement with Segmentation

Part (a): Using GrabCut for Segmentation

```

1 mask = np.zeros(image.shape[:2], np.uint8) # Initialize mask for GrabCut
2 rect = (65, 50, image.shape[1]-50, image.shape[0]-300) # Define the rectangle parameters that include the
3 bgdModel = np.zeros((1, 65), np.float64) # Create temporary arrays for the algorithm
4 fgdModel = np.zeros((1, 65), np.float64)
5
6 cv2.grabCut(image, mask, rect, bgdModel, fgdModel, 10, cv2.GC_INIT_WITH_RECT) # Apply GrabCut

```

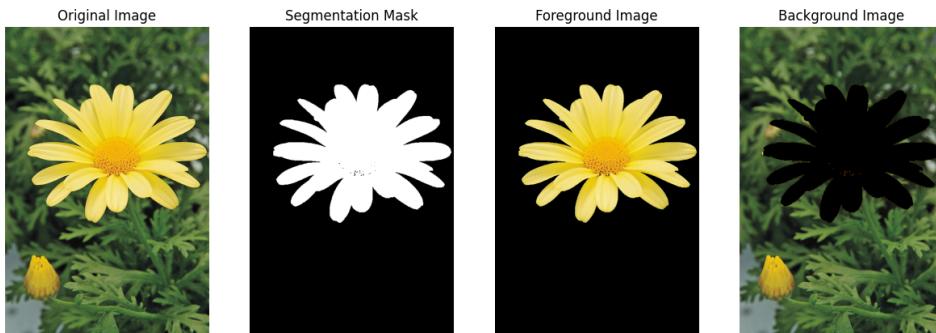


Figure 16: Original image, final segmentation mask, foreground image, and background image.

Part (b): Enhancing the Image

```

1 background = cv2.bitwise_and(image_rgb, image_rgb, mask=255-mask2) # Create the background
2 blurred_background = cv2.GaussianBlur(background, (27, 27), 0)
3 alpha = 0.5; beta = 0.5; # Combine the blurred background with the original foreground
4 enhanced_image = cv2.addWeighted(foreground, alpha, blurred_background, beta, 0)

```

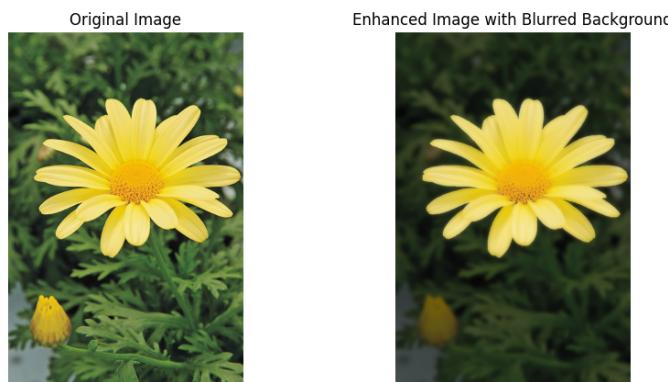


Figure 17: Original image(left) and Enhanced image with a blurred background(right).

Part (c): Discussing Background Darkness

The darkness around the flower edges in the enhanced image results from the blurring process. Blurring merges darker edge shadows into the lighter background, creating a contrast effect that accentuates the shadows, making the boundary appear darker than in the original.